


Matched distributed throughput

578 aligned post-warmup reports · median 21.84k → 24.29k tokens/s/GPU · 1.112×



Dashed lines are medians. Scheduled-save and input-stall windows are retained; no outliers are removed.