

Llama 3.1 8B on one GB200

B1 · S4096 · 32 layers · compiled standard cross-entropy · 31 timed steps

1.054x faster step · 5.1% less time

