


Working FP8 P/V training routes

Llama 3.1 8B · S4096 · common observed horizon 55.5B tokens · 1B-token EWM



Thin lines are downsampled raw loss; thick lines are causal token-domain EWMs.
The BF16 run is a historical, non-paired reference.